


Understanding Dispatching Approaches in the Scientific Python Ecosystem

- By Sebastian Berg and Aditi Juneja

Aditi Juneja (@Schefflera-Arboricola)

- **Grants/programs:** Google Summer of Code 2024, NumFOCUS's Small Development Grant (R3, 2024), CZI*
- **Open source projects/contributions:** nx-parallel, NetworkX (Core developer), scikit-image (dispatching), other scientific Python projects
- **Communities and Volunteering roles:** SciPy India (core-organiser), GSoC 2025 mentor (NetworkX), SciPy 2025 (proceedings paper reviewer), Scientific Python (maintainer - dispatch team), IndiaFOSS 2025 ("FOSS in Science" devroom manager), PyDelhi (reviewer), PyData Global 2024 (reviewer), PyCon US 2025 (proposal mentor)



Find work blogs, talk slides, CV, etc. at <https://github.com/Schefflera-Arboricola/blogs>

Contents

- Introduction to dispatching - 5 mins
- Dispatching in NumPy and for arrays - 6-7 mins
- Dispatching in networkx and for graphs - 10-12 mins
 - Scikit-image dispatching
- Overview - 4-5 mins
 - Spatch
 - Summary

Dispatching??

Arriving Sunday

Dispatched



Apsara Dustless Chalks | 4x Longer Than Regular Chalks | Hypoallergenic Chalk for Safe Using | Non-dust Chalk for Clean Writing | Available in Vibrant Colors | Ideal for Schools Box of 100 Chalks.

Sold by: Cocoblu Retail

₹343.00

Track package

Request cancellation

Return or replace items

Share gift receipt

Leave seller feedback

Write a product review

Dispatching– in general

Arriving Sunday

Dispatched



Apsara Dustless Chalks | 4x Longer Than Regular Chalks | Hypoallergenic Chalk for Safe Using | Non-dust Chalk for Clean Writing | Available in Vibrant Colors | Ideal for Schools Box of 100 Chalks.

Sold by: Cocoblu Retail

₹343.00

Track package


Request cancellation

Return or replace items

Share gift receipt

Leave seller feedback

Write a product review



Dispatching – in programming

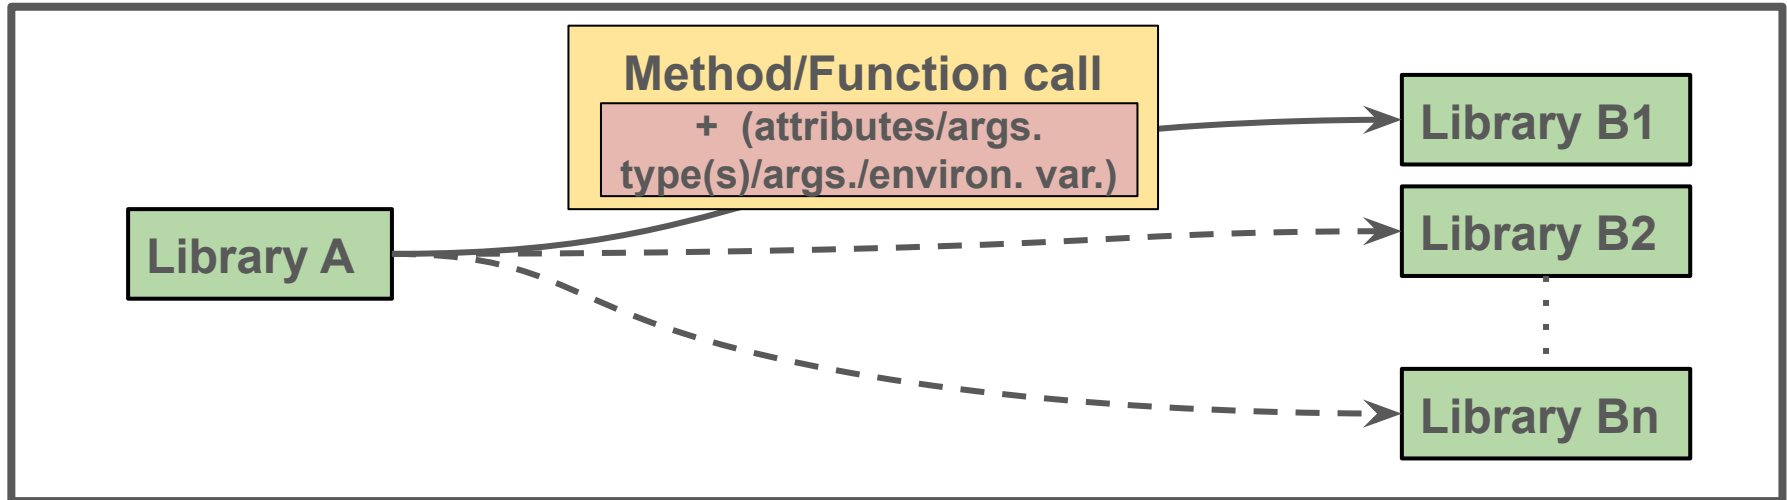
“In **computer science**, **dynamic dispatch** is the process of selecting which implementation of a **polymorphic** operation (**method** or function) to call at **run time**.”


Ref. https://en.wikipedia.org/wiki/Dynamic_dispatch

Dispatching – in programming

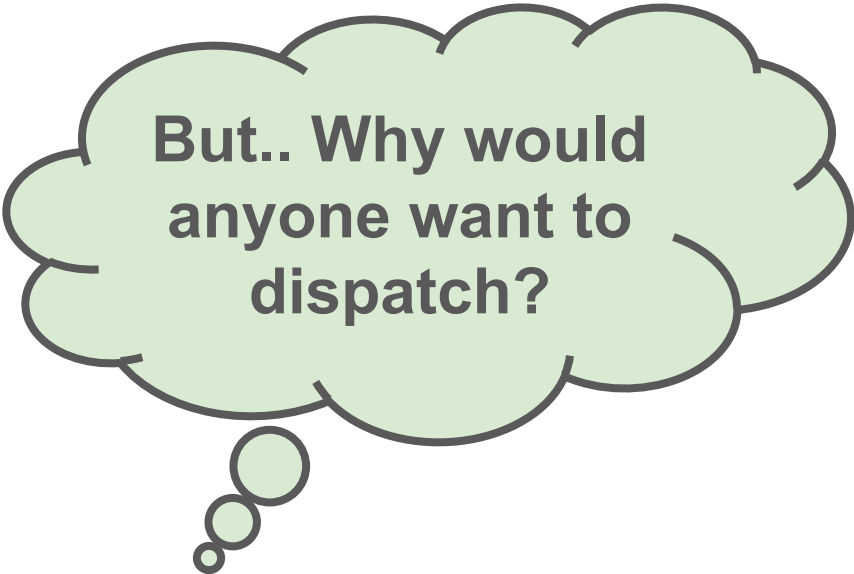
“In **computer science**, **dynamic dispatch** is the process of selecting which implementation of a **polymorphic** operation (**method** or function) to call at **run time**.”

Ref. https://en.wikipedia.org/wiki/Dynamic_dispatch





**But.. Why would
anyone want to
dispatch?**




**But.. Why would
anyone want to
dispatch?**

**Improve
Performance**

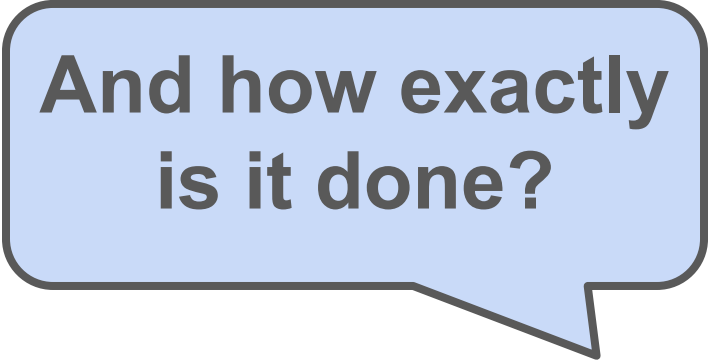
**Type
Compatibility**

**Simplify User
API**

**Enable different
workflows**



**But.. Why would
anyone want to
dispatch?**



**And how exactly
is it done?**

Example

What does dispatching mean here?

Inside ``library_A``

```
def add(x, y, z):  
    return x + y + z
```

Importing and Calling (hard-coded)

Inside `library_A`

```
def add(x, y, z, libB=None):  
    if libB == "library_B1":  
        import library_B1 as libB1  
        z = float(z)  
        return libB1.mod1.add(x, y, z)  
    elif libB == "library_B2" :  
        ...  
        ...  
    else:  
        return x + y + z
```

User

```
libA.add(1, 2, 3,  
libB="library_B1")
```

library_A

```
libB1.mod1.add(1, 2, 3.0)
```

library_B1

Steps involved

Inside `library_A`

```
def add(x, y, z, libB=None):  
    if libB == "library_B1":  
        import library_B1 as libB1  
        z = float(z)  
        return libB1.mod1.add(x, y, z)  
    elif libB == "library_B2":  
        ...  
        ...  
    else:  
        return x + y + z
```

Get `libB`

Import `libB`

Convert args for `libB`

Find `add` in `libB` and call it with converted args

Steps involved

Inside `library_A`

```
def add(x, y, z, libB=None):  
    if libB == "library_B1":  
        import library_B1 as libB1  
        z = float(z)  
        return libB1.mod1.add(x, y, z)  
    elif libB == "library_B2" :  
        ...  
        ...  
    else:  
        return x + y + z
```

Can we make
this more
general?

Steps involved

Inside `library_A`

```
def add(x, y, z, libB=None):  
    if libB == "library_B1":  
        import library_B1 as libB1  
        z = float(z)  
        return libB1.mod1.add(x, y, z)  
    elif libB == "library_B2" :  
        ...  
        ...  
    else:  
        return x + y + z
```

Can we make
this more
general?

... maybe we can
delegate some of these
steps to `libB`...

Step 4: Find `add` in `libB` and call it

Inside `library_A`

```
def add(x, y, z, libB=None):  
    if libB == "library_B1":  
        import library_B1 as libB1  
        z = float(z)  
        return libB1.mod1.add(x, y, z)  
    elif libB == "library_B2" :  
        ...  
        ...  
    else:  
        return x + y + z
```



Find `add` in `libB` and call it with converted args

Step 4: Find `add` in `libB` and call it

Inside `library_A`

```
def add(x, y, z, libB=None):  
    if libB == "library_B1":  
        import library_B1 as libB1  
        z = float(z)  
        all_funcs = libB1.get_all_funcs()  
        libb1_add = all_funcs.add  
        return libB1_add(x, y, z)  
    elif libB == "library_B2" :  
        ...  
        ...  
    else:  
        return x + y + z
```

Get a namespace* of all
the functions in `libB1`

Extract `add` from this
namespace

Call the extracted `add`

**Assumption: all functions in the
namespace have unique names*

Step 3: convert args for `add` in `libB`

Inside `library_A`

```
def add(x, y, z, libB=None):  
    if libB == "library_B1":  
        import library_B1 as libB1  
        z = float(z)  
        all_funcs = libB1.get_all_funcs()  
        libb1_add = all_funcs.add  
        return libB1_add(x, y, z)  
    elif libB == "library_B2" :  
        ...  
        ...  
    else:  
        return x + y + z
```

Convert args for `libB`



Step 3: convert args for `add` in `libB`

Inside `library_A`

```
def add(x, y, z, libB=None):  
    if libB == "library_B1":  
        import library_B1 as libB1  
        x, y, z = libB1.convert_args(add, x, y, z)  
        all_funcs = libB1.get_all_funcs()  
        libb1_add = all_funcs.add  
        return libB1_add(x, y, z)  
    elif libB == "library_B2" :  
        ...  
        ...  
    else:  
        return x + y + z
```

Convert function
in `libB`



Step 2: import `libB`

Inside `library_A`

```
def add(x, y, z, libB=None):  
    if libB == "library_B1":  
        import library_B1 as libB1  
        x, y, z = libb1.convert_args(add, x, y, z)  
        all_funcs = libB1.get_all_funcs()  
        libb1_add = all_funcs.add  
        return libB1_add(x, y, z)  
    elif libB == "library_B2" :  
        ...  
        ...  
    else:  
        return x + y + z
```

Import `libB`



Step 2: import `libB`

Inside `library_A`

```
def add(x, y, z, libB=None):  
    lib = __import__(libB)  
    if libB == "library_B1":  
        x, y, z = lib.convert_args(add, x, y, z)  
        all_funcs = lib.get_all_funcs()  
        lib_add = all_funcs.add  
        return lib_add(x, y, z)  
    elif libB == "library_B2" :  
        ...  
        ...  
    else:  
        return x + y + z
```

Generalising importing



Step 2: import `libB`

Inside `library_A`

```
def add(x, y, z, libB=None):  
    lib = __import__(libB)  
    if libB == "library_B1":  
        x, y, z = lib.convert_args(add, x, y, z)  
        all_funcs = lib.get_all_funcs()  
        lib_add = all_funcs.add  
        return lib_add(x, y, z)  
    elif libB == "library_B2" :  
        ...  
        ...  
    else:  
        return x + y + z
```

Generalising importing

Implication:
code inside
if-else also
became
generalised

Step 2: import `libB`

Inside `library_A`

```
def add(x, y, z, libB=None):  
    if libB != None:  
        lib = __import__(libB)  
        x, y, z = lib.convert_args(add, x, y, z)  
        all_funcs = lib.get_all_funcs()  
        lib_add = all_funcs.add  
        return lib_add(x, y, z)  
    else:  
        return x + y + z
```

Generalising importing

Getting rid of
`if-else`
conditions

Dispatching

Inside `library_A`

```
def add(x, y, z, libB=None):  
    if libB == "library_B1":  
        import library_B1 as libB1  
        z = float(z)  
        return libB1.mod1.add(x, y, z)  
    elif libB == "library_B2" :  
        ...  
        ...  
    else:  
        return x + y + z
```

Get `libB`

`library_A`

Import `libB`

`library_A`

Convert args
for `libB`

`library_B1`

Find `add` in `libB` and call
it with converted args

`library_B1`

Dispatching

Inside `library_A` (import and call)

```
def add(x, y, z, libB=None):
    if libB == "library_B1":
        import library_B1 as libB1
        z = float(z)
        return libB1.mod1.add(x, y, z)
    elif libB == "library_B2" :
        ...
        ...
    else:
        return x + y + z
```

Inside `library_A` (dispatching)

```
def add(x, y, z, libB=None):
    if libB != None:
        lib = __import__(libB)
        x, y, z = lib.convert_args(add,
x, y, z)
        all_funcs = lib.get_all_funcs()
        lib_add = all_funcs.add
        return lib_add(x, y, z)
    else:
        return x + y + z
```

Dispatching

Inside `library_A` (import and call)

```
def add(x, y, z, libB=None):
    if libB == "library_B1":
        import library_B1 as libB1
        z = float(z)
        return libB1.mod1.add(x, y, z)
    elif libB == "library_B2" :
        ...
        ...
    else:
        return x + y + z
```

Can we do something more here?

Inside `library_A` (dispatching)

```
def add(x, y, z, libB=None):
    if libB != None:
        lib = __import__(libB)
        x, y, z =
        lib.convert_args(add, x, y, z)
        all_funcs =
        lib.get_all_funcs()
        lib_add = all_funcs.add
        return lib_add(x, y, z)
    else:
        return x + y + z
```

Dispatching


Yes! use decorator.
Generalising for all functions

Inside `library_A`

```
@_dispatchable
def add(x, y, z):
    return x + y + z
```

```
def _dispatchable():
    @functools.wraps(func)
    def wrapper(*args, **kwargs):
        # check for libB kwarg in the function signature
        libB = kwargs.get("libB")
        try:
            lib = __import__(libB)
            args = lib.convert_args(func, *args, **kwargs)
            all_funcs = lib.get_all_funcs()
            lib_func = all_funcs.func
            return lib_func(args)
        except ImportError:
            return func(*args, **kwargs)
    return wrapper
```

**there are other ways
of dispatching as well.*




**How is dispatching
done in real projects?**

Dispatching in NumPy and for Arrays

Dispatching in NumPy

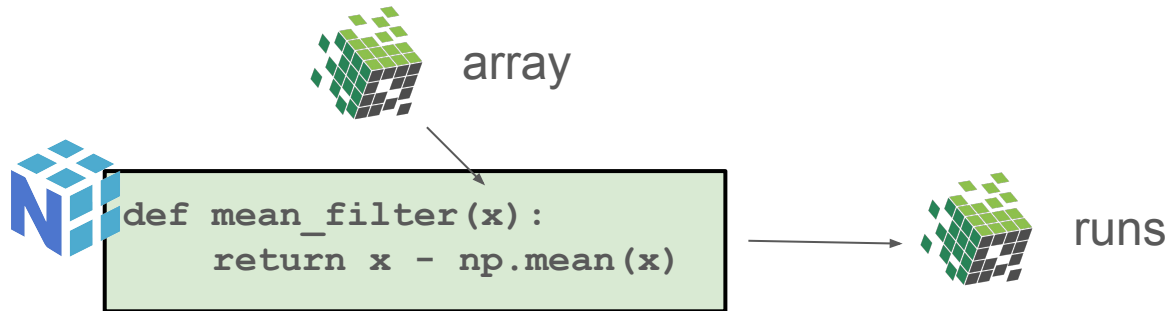
- Started around ~2013:
 - `__array_ufunc__` (NEP 13)
 - subset (“universal functions”)
 - `__array_function__` (NEP 18 + 35)
 - almost all other functions.
- Use-cases/users:
 - SciPy sparse
 - `astropy.units` (“NumPy array of meters”)
 - Dask array (distributed)
 - CuPy (GPU), JAX, cupynumeric...



Dispatching in NumPy – Usage

- Astropy quantities:
 - User use: ``np.mean`` not ``astropy.unit.mean``
- Cupy, etc.
 - `import cupy as np` possible
 - Dispatching allows function re-use!

Example:



Dispatching in NumPy – Implementation

`np.mean(cupy_array)`

for array args

Possible “backends”:

`hasattr(type(cupy_array), "__array_function__")`

For all unique types
(subclasses first)

Code in CuPy (simplified):

```
class ndarray:
    def __array_function__(self, func, types, *args, **kw):
        if not known_function(func): # np.mean known
            return NotImplemented

        for t in types:
            if not issubclass(t, (ndarray, numpy.ndarray)):
                return NotImplemented

        return cupy.mean(*args, **kw)
```

Can run?

Implementation

Dispatching in NumPy – Summary

- Allows interoperating with NumPy
- `__special_method__()` much like Python binary operators
- Adopted by many NumPy like objects
- Good approach if you own type

But:

- E.g. `torch` did not adopt (backwards compatibility concerns).
- Libraries (SciPy, etc.) do not use it (to allow cupy, JAX, arrays...)

Beyond NumPy – Array API

- NEP 37 (not implemented) evolved into **Array API**
 - `xp = array_api_compat.array_namespace(*arrays)`
 - `xp` is `numpy`, `cupy`, `torch`, `jax.numpy`...
 - `xp` *explicit* state: dispatch only once
 - **Being adopted** into **SciPy**, **sklearn**, ...
- For dataframes: **narwhals**


Clear **library focus**. End-users: `import cupy as np` is more practical.
(but every end-user also writes their small functions/libraries)

Dispatching in NetworkX for Graphs

What is networkx?


**NetworkX is a
graph analysis
library**

What are graphs/networks?



Need for dispatching in NetworkX

- **Speed:**
 - NetworkX is pure python, therefore slow;
 - but simple to understand, use, maintain, and contribute to
- **Standardisation** - Simplifying user API



Steps involved in dispatching

Inside `library_A`

```
def add(x, y, z, libB=None):  
    if libB == "library_B1":  
        import library_B1 as libB1  
        z = float(z)  
        return libB1.mod1.add(x, y, z)  
    elif libB == "library_B2" :  
        ...  
        ...  
    else:  
        return x + y + z
```

Get `libB`

Import `libB`

*

Convert args for `libB`

Find `add` in `libB` and call it with converted args

Step 1: Get the backend package name from the user

Step 1: Get the backend package name from the user

Speed-up/usage **demo** with cugraphs,
graphblas and/or parallel backends

<https://colab.research.google.com/drive/16bUxGvBoBAg1dBRc6yfxixlo7DocM8s?usp=sharing>

4 ways of to get backend name in NetworkX

```
nx.betweenness centrality(CuG)
```

CuG.__networkx_backend__
contains the backend name
as a string

```
nx.betweenness centrality(G,  
backend = "parallel")
```

```
with nx.config(backend_priority =  
["cugraph", "graphblas", "parallel"]):  
    nx.betweenness centrality(G)
```

```
$ NETWORKX_BACKEND_PRIORITY="graphblas"  
$ python nx_code.py
```

Inside `networkx`

```
@_dispatchable  
def betweenness centrality(  
    G, k, ... ,seed  
):  
    ...
```

nx-parallel

nx-cugraph

Python-
graphblas

One-to-one mapping of
backends and types

Step 2: Getting the function from the backend package

Step 2: Getting the function from the backend package

Python entry-points to get all the installed backends (and their metadata – supported functions and convert functions).

What are entry-points?

Entry-points are used to extend a functionality of a library.

What are entry-points?

Entry-points are used to extend a functionality of a library.

Main library



Entry point

Extension (or plug-in) libraries



Dispatching with entry-points is exploiting entry-points to the max—

i.e. change the whole machinery inside the vacuum cleaner(main library's implementation)— and just keep the outside red UI-plastic-buttons frame (user-API).

Defining an entry-point in the backend library

In backend's `pyproject.toml`

```
[project.entry-points."networkx.backends"]
```

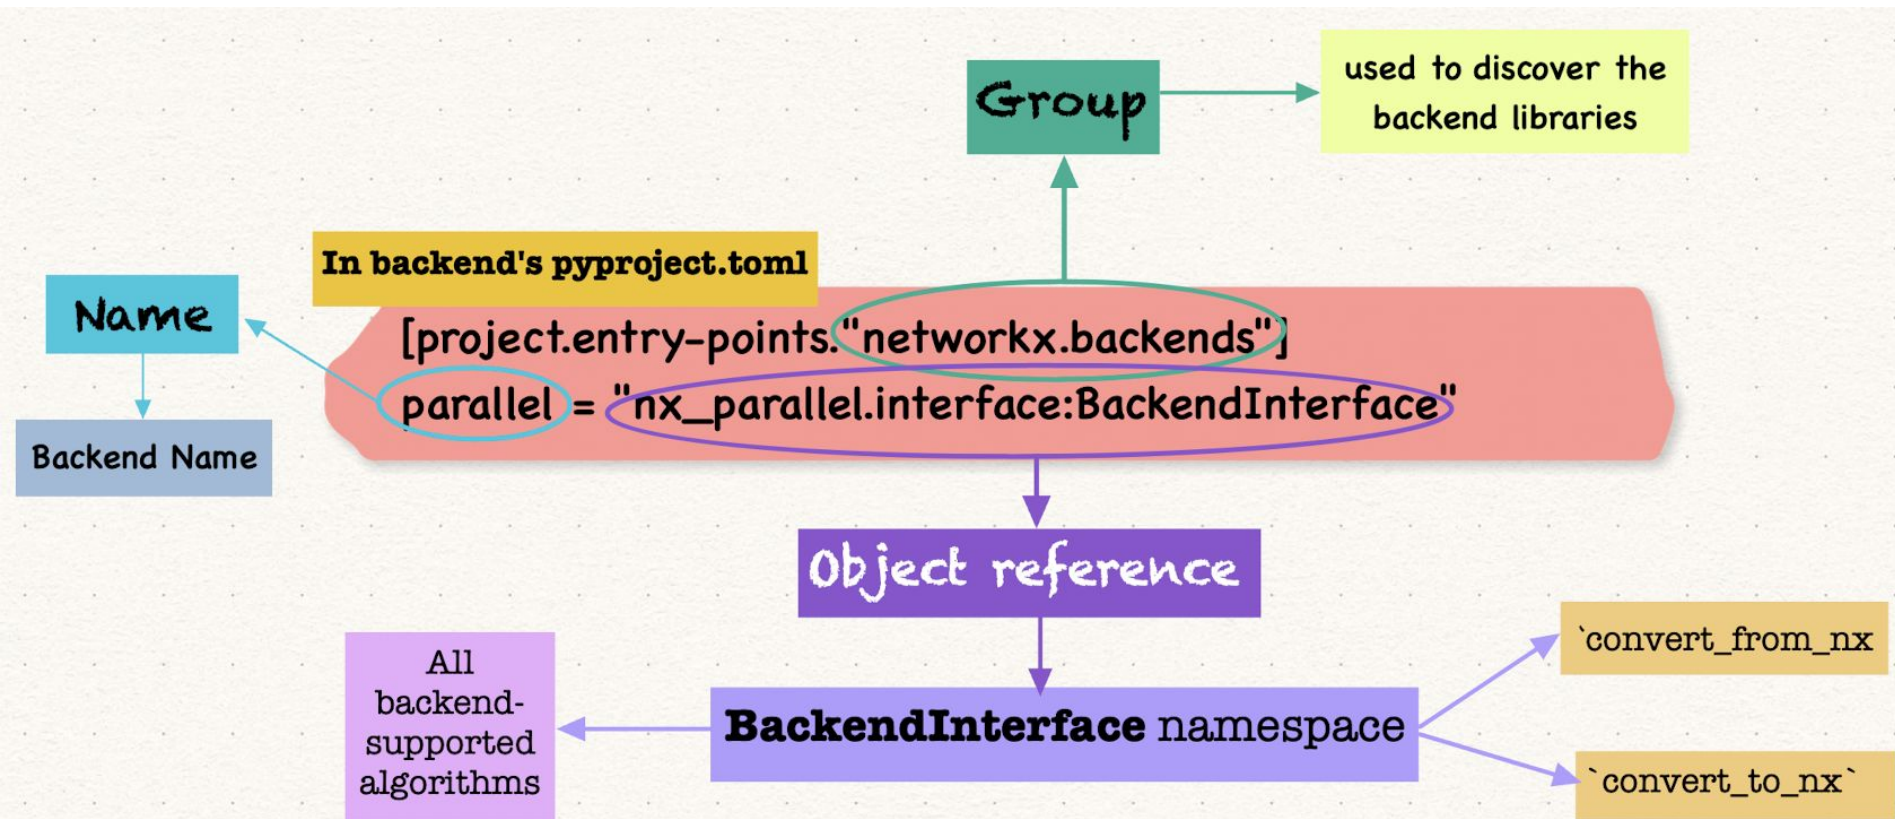
```
parallel = "nx_parallel.interface:BackendInterface"
```

Group

Name

Object reference

Defining an entry-point in the backend library



How does entry-point based dispatching work?

**Grouping all the
installed backends**

```
eps = entry_points(...)
```

How does entry-point based dispatching work?

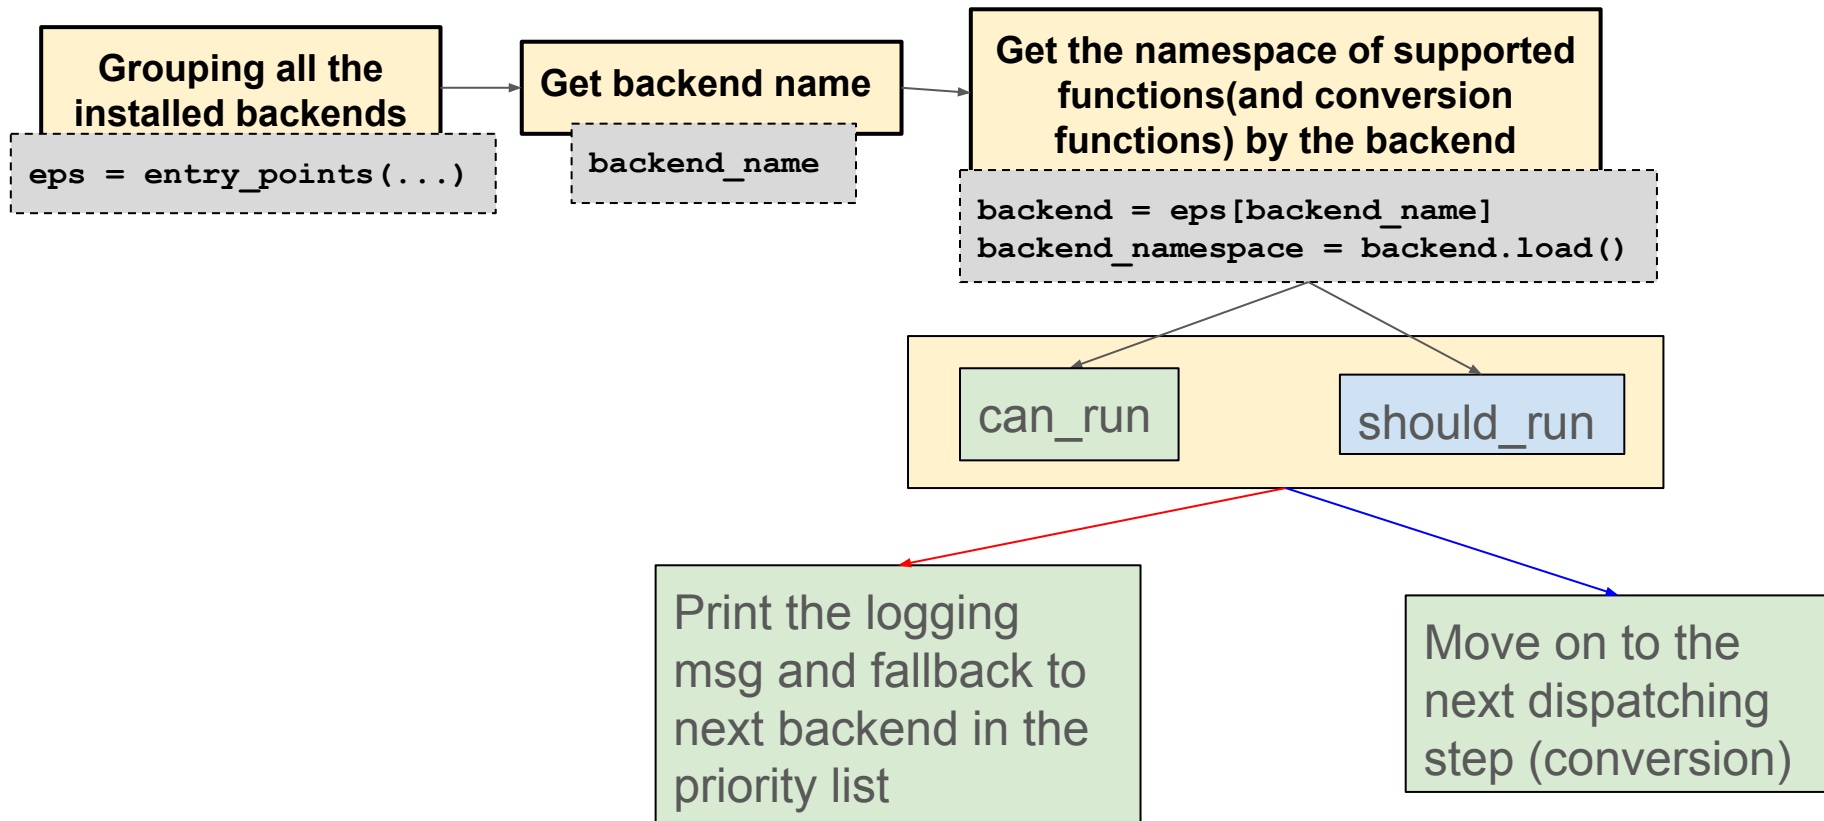
**Grouping all the
installed backends**

```
eps = entry_points(...)
```

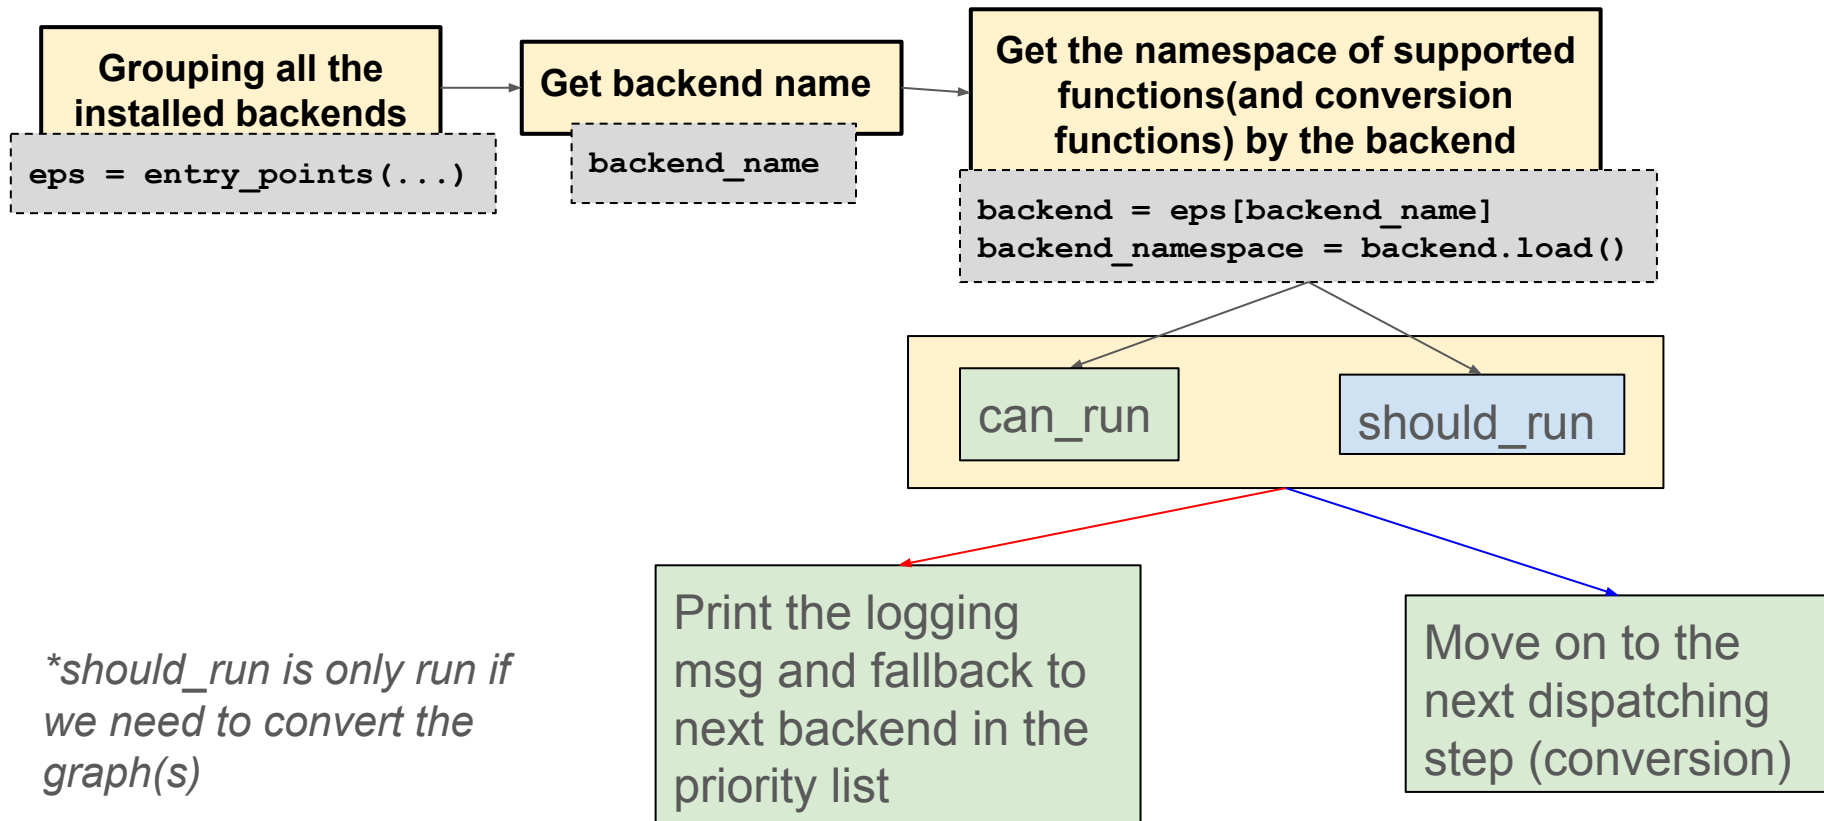
Get backend name

```
backend_name
```

How does entry-point based dispatching work?



How does entry-point based dispatching work?



How does entry-point based dispatching work?

Grouping all the
installed backends

```
eps = entry_points(...)
```

Get backend name

```
backend_name
```

Get the namespace of supported
functions(and conversion
functions) by the backend

```
backend = eps[backend_name]  
backend_namespace = backend.load()
```

Conversion step, if
backend-name-based)

```
backend_namespace.convert_from_nx(*args, **kwargs)
```


How does entry-point based dispatching work?

Grouping all the installed backends

```
eps = entry_points(...)
```

Get backend name

```
backend_name
```

Get the namespace of supported functions (and conversion functions) by the backend

```
backend = eps[backend_name]  
backend_namespace = backend.load()
```

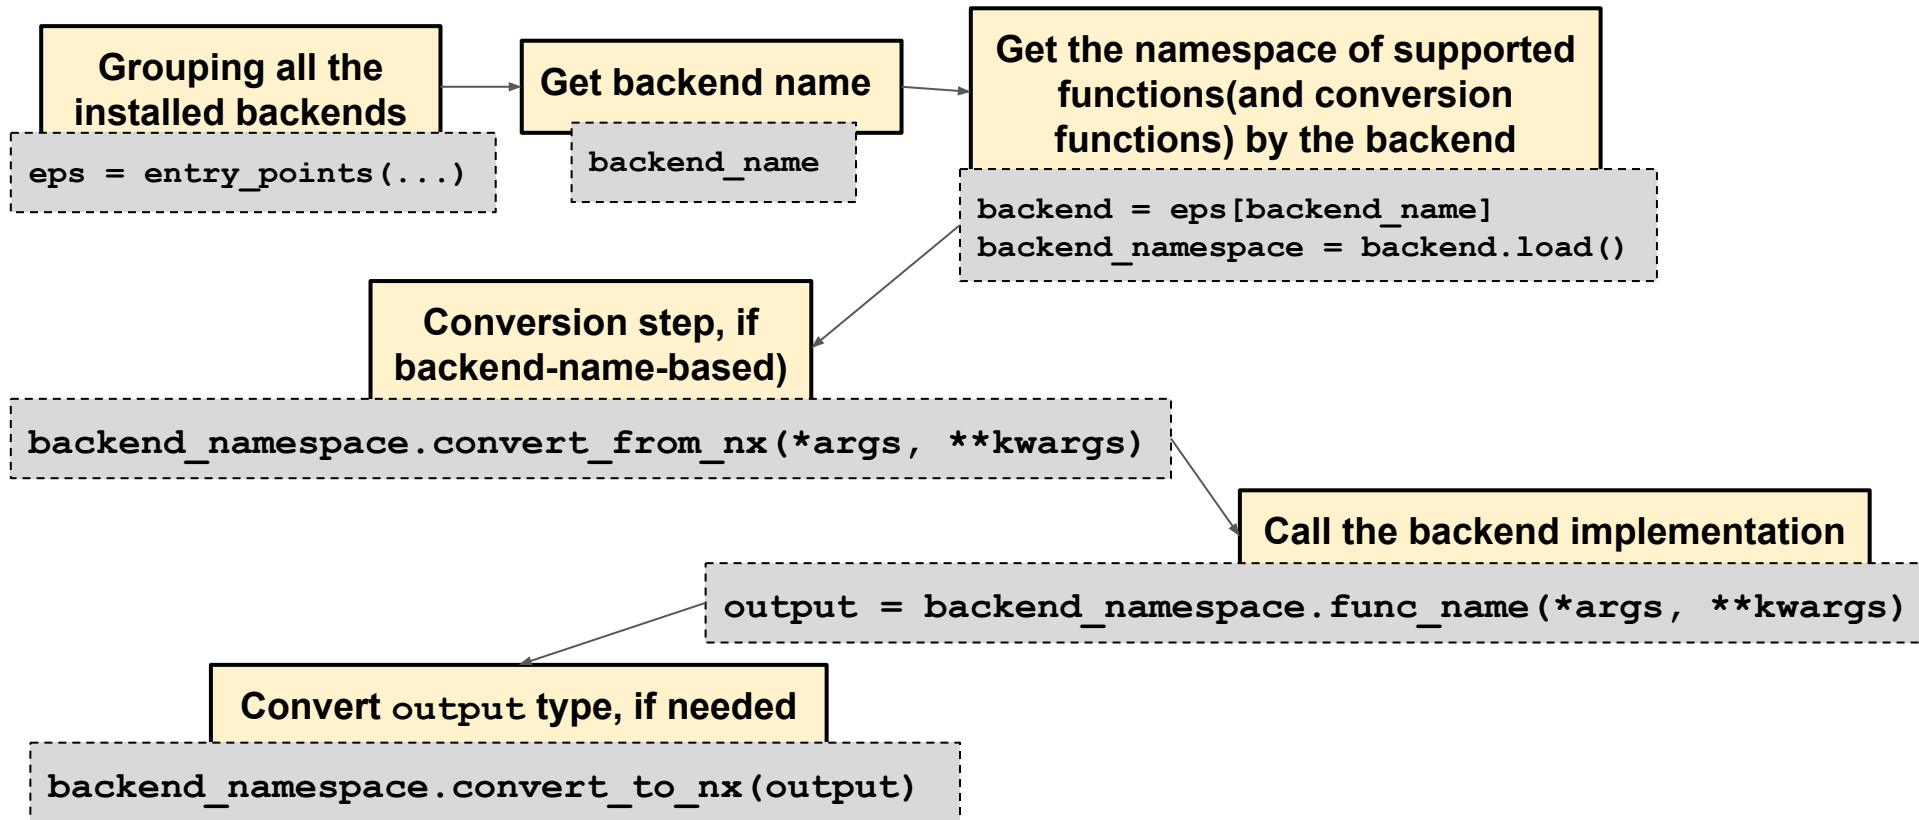
Conversion step, if backend-name-based)

```
backend_namespace.convert_from_nx(*args, **kwargs)
```

Call the backend implementation

```
output = backend_namespace.func_name(*args, **kwargs)
```

How does entry-point based dispatching work?



Other features of NetworkX dispatching (if time permits)

- Backend testing
- Docs : second entry point
- By looking at the `.backends` attribute, you can get the set of all currently installed backends that implement a particular function. For example:

```
>>> nx.betweenness centrality.backends
```

```
{'parallel'}
```

- Specialised backend priority for algorithms, generators,.. etc.
- Logging
- Fallback

Read more at <https://networkx.org/documentation/latest/reference/backends.html>

Dispatching in scikit-image

- Image object (`numpy.ndarray`) → not owned by scikit-image
 - cannot add backend name to the array object (like `__networkx_backend__` in NetworkX)

Dispatching in scikit-image

- Image object (`numpy.ndarray`) → not owned by scikit-image
 - cannot add backend name to the array object (like `__networkx_backend__` in NetworkX)
- **Goal:** want it to look like type based dispatching (with a component of backend selection) – internally entry-point based.

Dispatching in scikit-image

- Image object (``numpy.ndarray``) → not owned by scikit-image
 - cannot add backend name to the array object (like `__networkx_backend__` in NetworkX)
- **Goal:** want it to look like type based dispatching (with a component of backend selection) – internally entry-point based.
- Why not use Array API standards? – Hard for Cython-heavy array consuming libraries to adopt these standards– therefore using entry-points for dispatching.

Dispatching in scikit-image

- Image object (``numpy.ndarray``) → not owned by scikit-image
 - cannot add backend name to the array object (like `__networkx_backend__` in NetworkX)
- **Goal:** want it to look like type based dispatching (with a component of backend selection) – internally entry-point based.
- Why not use Array API standards? – Hard for Cython-heavy array consuming libraries to adopt these standards– therefore using entry-points for dispatching.
- **Current state:** No array conversions in dispatching code! – backend developers and users need to take care of the types.

Dispatching in scikit-image

- Image object (``numpy.ndarray``) → not owned by scikit-image
 - cannot add backend name to the array object (like `__networkx_backend__` in NetworkX)
- **Goal:** want it to look like type based dispatching (with a component of backend selection) – internally entry-point based.
- Why not use Array API standards? – Hard for Cython-heavy array consuming libraries to adopt these standards– therefore using entry-points for dispatching.
- **Current state:** No array conversions in dispatching code! – backend developers and users need to take care of the types.

Possible Solution : spatch

Spatch – <https://github.com/scientific-python/spatch>

Spatch – <https://github.com/scientific-python/spatch>

Create library or blueprint for the type of dispatching scikit-image needs!
(based on NetworkX/skimage lessons.)

Goal:

- Help implement good dispatching and backends
- Develop a recognizable API for users

Principles:

- Entry-points (docs + introspection)
- Fast *type dispatching* (caching)
- *Named backend selection* support

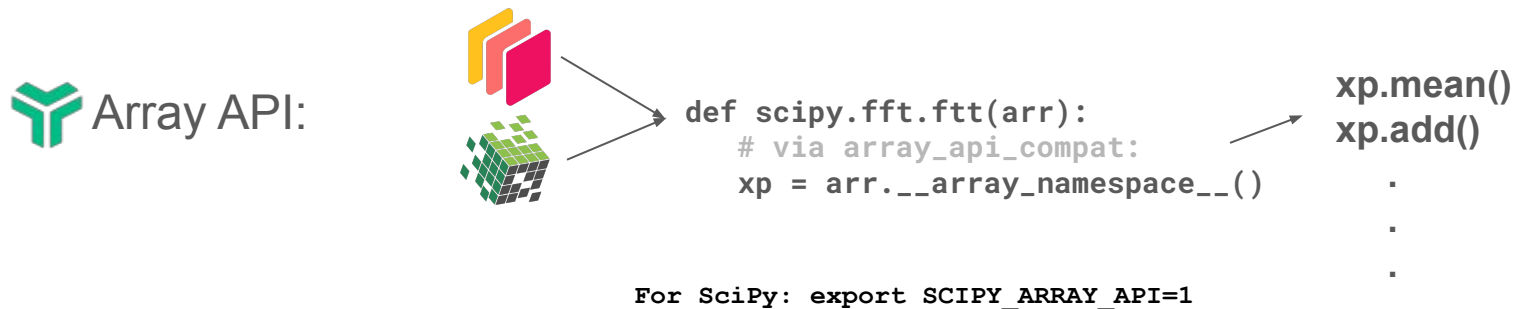

Spatch – Key points

- Library:
 - Mostly add decorator.
- Backends – entrypoint:
 - Which types: "`~cupy:ndarray`" (cupy arrays and subclasses)
 - Auto-generated mapping:

```
functions = {  
    "skimage.filters.gaussian": {  
        "impl": "cucim.skimage.filters.gaussian"}  
}
```
- User API:
 - (implicit) type dispatching
 - `with backend_opts(prioritize="opencv")`
 - ...

Summary

Summary – Type based



Summary – Type + Selection



`nx.betweenness centrality(G)`



Backend name from graph object:

- `G.__networkx_backend__`

Backend name directly via user:

- **`NETWORKX_BACKEND_PRIORITY`**
- ``nx.config()``
- ``backend=``



entry-points



nx-parallel



nx-cugraph



GRAPHBLAS



with spatch

Via object:

- `"cupy:ndarray", ...`

User:

- environment / `with config`
- (future)



entry-points



Overview

	NumPy	Array API	NetworkX	spatch
Type dispatching	<code>__array_function__()</code>	<code>__array_namespace__()</code>	<code>__networkx_backend__</code>	"module:name"
Selection (speed)	no	no	backend="name" env / with config()	with backend_opts() / env
entry-point	no	no	yes	yes
conversion	by backend	no	If needed (by convert functions)	by backend
Focus	end-user	Array consuming library	End-user (speed)	Help devs dispatch for end-users

Overview

	NumPy	Array API	NetworkX	spatch
Type dispatching	<code>__array_function__()</code>	<code>__array_namespace__()</code>	<code>__networkx_backend__</code>	"module:name"
Selection (speed)	no	no	backend="name" env / with config()	with backend_opts() / env
entry-point	no	no	yes	yes
conversion	by backend	no	If needed (by convert functions)	by backend
Focus	end-user	Array consuming library	End-user (speed)	Help devs dispatch for end-users

Thank you!

Dispatching in Scientific Python ecosystem

- SPEC 2 : <https://scientific-python.org/specs/spec-0002/>
- spatch : <https://github.com/scientific-python/spatch/>
- Array API standards: <https://data-apis.org/array-api/latest/>
- NumPy's type-based dispatching
 - <https://numpy.org/neps/nep-0037-array-module.html>
 - <https://numpy.org/neps/nep-0047-array-api-standard.html>
- NetworkX
 - <https://networkx.org/documentation/latest/reference/backends.html>
 - <https://networkx.org/documentation/latest/reference/configs.html>
 - Dispatch meetings: <https://scientific-python.org/calendars/networkx.ics>
 - <https://github.com/networkx/networkx/issues?q=is%3Aissue%20state%3Aopen%20label%3ADispatching>
- Scikit-image
 - [scikit-image-PR#7520](https://github.com/scikit-image/scikit-image/pull/7727)
 - <https://github.com/scikit-image/scikit-image/pull/7727>
 - <https://github.com/rapidsai/cucim/issues/829>
- Scikit-learn
 - <https://github.com/scikit-learn/scikit-learn/pull/30250>
 - <https://youtu.be/f42C1daBNrg?si=A9mZ2mZd2HzEhu8S>
- SciPy's Array API adoption
 - https://docs.scipy.org/doc/scipy/dev/api-dev/array_api.html
 - https://youtu.be/16rB-fosAWw?si=ys_-ZTnUKvO_aZKu
- DataFrame API standards
 - <https://github.com/narwhals-dev/narwhals>
 - <https://data-apis.org/dataframe-api/draft/>
- Scientific Python discord(#dispatching thread): <https://discord.com/invite/vur45CbWmZ>

FIN.

